



VECTORIZATION (SIMD)

Alessandro Casalino
[a.casalino@cineca.it](mailto:a.casalino@ Cineca.it)

July 6, 2023

CINECA

OUTLINE

- ▶ Introduction to vectorization
- ▶ Auto-vectorization
- ▶ Vectorization in OpenMP

Introduction


WHY VECTORIZATION?

```
void add(double* a, double* b,  
         double* c)  
{  
    for(size_t i = 0; i < 16; i++) {  
        a[i] = b[i] + c[i];  
    }  
}
```

- ▶ Each iteration is done sequentially one by one
- ▶ Even if parallelized, each thread do one assigned iteration one by one

Can we do better?

SCALAR INSTRUCTION




Each iteration is done sequentially one by one.

Operations

- ▶ 16 + 16 memory loads = 32 memory loads
- ▶ 16 additions
- ▶ 16 store

VECTORIZED INSTRUCTION



Operations are performed on a block of data (**vector**).

Operations

- ▶ 4 + 4 memory loads = 8 memory loads
- ▶ 4 additions
- ▶ 4 store

SIMD (SINGLE INSTRUCTION MULTIPLE DATA)


SIMD register



↑
Lane

- ▶ A **SIMD Register** holds many values of the same datatype
- ▶ Each value is a **SIMD lane**
- ▶ SIMD hardware instructions modify the vector registers
- ▶ SIMD instructions can operate on all (or part) of the values in a SIMD register
- ▶ Width in recent CPUs up to 512 bit

WHY VECTORIZATION?



- ▶ OpenMP vectorization = SIMD
- ▶ Operates at once on an entire block, i.e. **operations on multiple data concurrently**
- ▶ **Dedicated registers**
- ▶ **Maximize bandwidth**

Auto-vectorization

AUTO-VECTORIZATION

Modern compilers **auto-vectorize** loops, provided that:

- ▶ complexity of the code is low enough for the compiler to select proper instructions;
- ▶ code pattern is recognized by the compiler;
- ▶ the loop benefits from a vectorization.

Compiler Reports might help

- ▶ `icc -qopt-report=N -qopt-report-phase=loop,vec`
- ▶ `gcc -ftree-vectorize -ftree-vectorize-verbose [-O3]`

Note

`icc` enables vectorization by default.
`gcc` enables vectorization with `-O3`.

VECTORIZE THE CODE

```
31 for (size_t i=0; i<n; ++i)
32     for (size_t k=0; k<n; k++)
33         for (size_t j=0; j<n; ++j)
34             c[i][j] += a[i][k]*b[k][j];
```

Note

Note that only the last loop is vectorized (line 33).

icc Compiler Report

```
LOOP BEGIN at mat_prod.c(33,3)
  remark #15300:  LOOP WAS VECTORIZED
  remark #15442:  entire loop may be executed in remainder
  remark #15448:  unmasked aligned unit stride loads:   2
  remark #15449:  unmasked aligned unit stride stores:  1
  remark #15475:  -- begin vector cost summary --
  remark #15476:  scalar cost:   9
  remark #15477:  vector cost:  4.000
  remark #15478:  estimated potential speedup:  2.210
  remark #15488:  -- end vector cost summary --
  remark #25015:  Estimate of max trip count of loop=125
LOOP END
```

AUTO-VECTORIZATION FAILING CAUSES

- ▶ **Data dependencies**
- ▶ **Alignment**
- ▶ **Function calls in the loop block**
- ▶ **Complex control flow/conditional branches**
- ▶ Loop not countable (e.g., upper bound not runtime constant)
- ▶ Mixed data types
- ▶ Non-unit stride between elements
- ▶ **Loop body too complex**
- ▶ Vectorization seems inefficient
- ▶ ...

DATA DEPENDENCY: FLOW DEPENDENCY

```
for(size_t i = 0; i < n; i++) {  
    a[i] = b[i] + 1;  
    c[i] = a[i] + 1;  
}
```

RAW (Read After Write)

All the values are available within the loop iteration: **vectorization possible**.

i = 0

a[0] = b[0] + 1;

c[0] = **a[0]** + 1;

i = 1

a[1] = b[1] + 1;

c[1] = **a[1]** + 1;

DATA DEPENDENCY: ANTI-DEPENDENCY

```
for(size_t i = 1; i < n; i++) {  
    a[i] = b[i] + 1;  
    c[i] = a[i-1] + 1;  
}
```

WAR (Write After Read)

A value should be loaded to a previous iteration: **vectorization not possible**.

i = 1

a[1] = b[1] + 1;

c[1] = **a[0]** + 1;

i = 2

a[2] = b[2] + 1;

c[2] = **a[1]** + 1;

DATA DEPENDENCY

Flow Dependency

```
for(size_t i = 0; i < n-1; i++) {  
    a[i] += a[i+1];  
}
```

icc Compiler Report

```
remark #15300:  LOOP WAS  
VECTORIZED
```

godbolt.org/z/q5G1WcYbn

Anti-Dependency

```
for(size_t i = 1; i < n; i++) {  
    a[i] += a[i-1];  
}
```

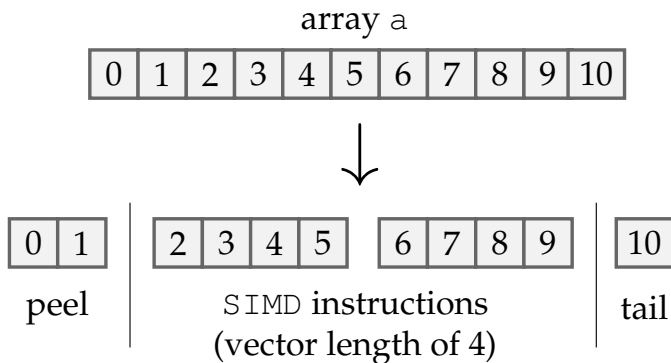
icc Compiler Report

```
remark #15344:  loop was not  
vectorized:  vector dependence  
prevents vectorization.
```

godbolt.org/z/azn5nEohM

DATA ALIGNMENT

```
for(size_t i = 0; i < 11; i++) {  
    a[i] = 1.;  
}
```



- ▶ For best effect: **starting address of vectors should be aligned on the correct boundary**
- ▶ Address in memory should be multiple of vector length in bytes

Otherwise **unaligned data**

- ▶ May require **multiple data loading, cache lines and instructions**
- ▶ Generates 3 different version of loops: **peel, kernel, tail**

DATA ALIGNMENT

```
31 for (size_t i=0; i<n; ++i)
32     for (size_t k=0; k<n; k++)
33         for (size_t j=0; j<n; ++j)
34             c[i][j] += a[i][k]*b[k][j];
```

icc Compiler Report

LOOP BEGIN at mat_prod.c(33,3)

<Remainder loop for vectorization>

remark #15335: remainder loop was not vectorized:
vectorization possible but seems inefficient. Use vector always
directive or -vec-threshold0 to override

remark #15442: entire loop may be executed in remainder

remark #15448: unmasked aligned unit stride loads: 2

remark #15449: unmasked aligned unit stride stores: 1

remark #15475: -- begin vector cost summary --

remark #15476: scalar cost: 9

remark #15477: vector cost: 4.000

remark #15478: estimated potential speedup: 2.210

remark #15488: -- end vector cost summary --

LOOP END

Vectorization in OpenMP

SIMD CONSTRUCT

Serial

```
#pragma omp simd
for(size_t i = 0; i < n; i++) {
    a[i] = b[i] + c[i];
}
```

- ▶ `simd` tells the compiler to vectorize
- ▶ No worksharing, i.e., only enables SIMD
- ▶ Compiler chooses the appropriate vector length for the target architecture

Worksharing

```
#pragma omp for simd
for(size_t i = 0; i < n; i++) {
    a[i] = b[i] + c[i];
}
```

- ▶ `for` tells the compiler to start the parallelization on multiple threads
- ▶ `simd` enables the vectorization for each thread
- ▶ First the work is shared between threads, then vectorization on chunks fitting the SIMD register

DATA SHARING

```
void foo (double * a, double * b,  
         double * c, int n) {  
    int i;  
    double t1, t2;  
  
    #pragma omp simd private(t1,t2)  
    for(i=0; i<n; i++) {  
        t1 = func1(b[i], c[i]);  
        t2 = func2(b[i], c[i]);  
        a[i] = t1 + t2;  
    }  
}
```

- ▶ Loop iterator `i` is private
- ▶ Memory pointed by `a`, `b` and `c` is shared

SIMD supported clauses

- ▶ `private(list)`
- ▶ `lastprivate(list)`
- ▶ `reduction(operation:list)`

SAFELLEN CLAUSE

```
#pragma omp simd safelen(4)
for(i=0; i<n-4; i++) {
    a[i] += a[i+4];
}
```

`safelen` specifies the **distance between iterations where it is safe to vectorize**.

- ▶ Useful to specify the maximum number of iterations that run concurrently without breaking a dependence.
- ▶ In practice, maximum SIMD vector length.

Best practice

Avoid using this clause. Specifying explicit vector lengths builds in obsolescence to the code as hardware vector lengths continually change.

COLLAPSE CLAUSE

`simd` partitions the loop iterations into chunks. But the partitioning requires only one iteration space to be valid. We can use `collapse` to solve this issue.

```
float A[4][4]
#pragma omp simd collapse(2)
for(size_t i=0; i < 4; i++) {
    for(size_t j=0; j<4; i++) {
        A[i][j] += 1;
    }
}
```

Best practice

This is the correct way to collapse for loops for `simd`.

```
float A[4][4]
#pragma omp simd
for(size_t k=0; k < 16; k++) {
    size_t i = k / 4;
    size_t j = k % 4;
    A[i][j] += 1;
}
```

Warning

Access to `A` is not linear with respect to the collapsed loop.

OTHER SIMD CLAUSES

`simdlen(length)`

- ▶ Specify the preferred length of SIMD registers
- ▶ Must be less or equal than `safelen`

Best practice

Again, since the hardware architectures changes, avoid using these *magic numbers* in the code.

`linear(length)`

- ▶ Specify the variable's value in relationship with the iteration number

Take-Home Messages

TAKE-HOME MESSAGES

- ▶ **OpenMP SIMD** constructs instructs the compiler to vectorize exploiting modern CPUs vector accelerator
- ▶ The **performance improvements with vectorization can be substantial**
- ▶ **Modern compilers are capable of auto-vectorize loops**, provided some conditions are met
- ▶ **Check the compiler reports to understand if the loops are correctly vectorized**

Exercises

EXERCISES



Matrix Multiplication

- ▶ Add `simd` constructs
- ▶ Run and time it

Value of π

- ▶ Add `simd` constructs
- ▶ Run and time it

Thank you!